MUZZAMIL RASULLY

Forward Deployed Engineer | Production AI Agents and Evaluation | Python, TypeScript, MCP
+44 7311 686838 | Rasully.m1@gmail.com | muzzamilrasully.com | github.com/Muzzamil-Codes-dev
London | Full right to work in the UK, no sponsorship required

EXPERIENCE

AI Consultant, Sola, early-stage AI consultancy — Aug 2026 - Present

- Built and deployed a Python/FastAPI document-processing application for a six-person operations team handling 1,200 documents a month, combining retrieval with schema-validated LLM extraction and cutting the time per document from 15 to 6 minutes (60%).
- Made it hold up in production: a background queue so failed jobs retry without blocking users, human review for outputs that fail validation, and failure monitoring on AWS (Linux, PostgreSQL). Worked with minimal supervision and flagged where AI was not the right answer.

AI Engineer, Ginseng — Aug 2025 - Aug 2026

- Built a 300-case evaluation suite covering missing fields, unsupported claims, malformed outputs and incorrect tool selection, and gated it in GitHub Actions (targeted checks on pull requests, the full suite before staging); it blocked 14 critical regressions, including 5 hallucinations, and cut production incidents by 40%.
- In a team of three AI engineers, built Action Sync with Cursor, an internal assistant used by around 200 staff: Jira integration via MCP connectors, a meeting agent that joins calls and returns structured notes, document summaries and a usage dashboard.
- Ran stakeholder discovery and shipped Python AI code on Grok models, integrated into enterprise systems across 6+ concurrent client workstreams, using tool calling, structured outputs and fallback handling over authorised, access-controlled client data, and documented behaviour and limits for each team.
- Built RAG pipelines on AWS over 100,000+ records, tuning accuracy, latency and cost, and cut manual processing time by 30%.

Founder, Sharpmind, online tutoring agency — Jul 2026 - Present

- Built, deployed and run an LLM agent end to end: it automates approximately 70% of routine tasks and cut replies from 2-3 hours to under 10 minutes; prompt and guardrail changes tested through A/B across 100+ live interactions cut manual intervention by a further approximately 60%.

Sales Manager, Flatpay, B2B fintech — Dec 2025 - Feb 2026

- Sold payment technology directly to business owners, learning each operation quickly and explaining a technical product and its trade-offs in plain language.

PROJECTS

Bug-triage agent service with evaluation and human review — Personal project

Stack: Python, FastAPI, Claude API, LangGraph, E2B, React, SQLite, pytest, Docker, OpenTelemetry

- Supervisor and specialist agents ingest unstructured bug reports, reproduce issues in an E2B sandbox, and return schema-validated triage through a REST API; low-confidence results escalate to a human, with a dead letter queue and OpenTelemetry tracing.
- Regression-evaluation harness with 88% severity accuracy on 25 labelled cases; a cache cut repeat runs from 13s to 0.5s, and a 36-test pytest suite covers 99% of the back end.

Retrieval benchmark, Natural Questions — Personal project

- Hybrid retrieval with cross-encoder reranking reached 93.80% Recall@1 across 871 queries against 54.76% for TF-IDF; I reported the failed query-expansion experiment rather than dropping it.

EDUCATION

BSc Mathematics, 2:1, Queen Mary University of London — Sept 2022 - May 2026

Relevant modules: Machine Learning; Natural Language Processing; Statistical Modelling; Probability; Linear Algebra.

TECHNICAL SKILLS

Agents and LLMs: Grok (1+ years building with it), Claude and OpenAI APIs, tool calling, MCP connectors, LangGraph, RAG, structured outputs, evaluation harnesses, tracing
Engineering: Python, TypeScript, SQL, FastAPI, REST APIs, React, PostgreSQL, Docker, Linux, AWS, Git, GitHub Actions CI/CD, pytest
AI-assisted development: Cursor, Claude and GitHub Copilot daily, reviewing generated code rather than accepting it