MUZZAMIL RASULLY

AI Engineer | Production LLM Applications | RAG, Agent Design and Orchestration | Python, AWS, CI/CD

+44 7311 686838 | Rasully.m1@gmail.com | muzzamilrasully.com | github.com/Muzzamil-Codes-dev

London, willing to relocate | Full right to work in the UK, no sponsorship required

EXPERIENCE

AI Consultant, Sola, early-stage AI consultancy — Aug 2026 - Present

- Design, build and ship production LLM applications in Python: retrieval-augmented generation over large document sets, and agent systems in LangGraph and LangChain where a supervisor agent validates specialist output against a schema before release.
- Design and refine agent behaviours, prompts and retrieval approaches, then review live performance and adjust, rather than shipping a prototype and walking away.
- Fine-tune models with LoRA and RLHF where retrieval and prompting are not the right answer, working with minimal supervision in a small cross-functional team and saying plainly when AI is not the right answer at all.

AI Engineer, Ginseng — Aug 2025 - Aug 2026

- Wrote, tested and deployed AI and LLM code in Python integrated with existing enterprise systems across 6+ concurrent workstreams, applying models from Anthropic and OpenAI to classification, extraction and summarisation with structured outputs, tool calling and fallback handling.
- Built RAG and data pipelines over 100,000+ records on AWS covering ingestion, chunking strategy, embeddings, vector search and reranking, tuned jointly for accuracy, latency and running cost, cutting manual processing time by 30%.
- Reviewed AI feature performance and monitoring data continuously: built ground-truth datasets and automated test harnesses, measured precision, recall, F1 and AUC, ran failure-mode analysis, and implemented drift monitoring that caught degradation before it reached users.
- Led a 10 person cross-functional initiative producing a reusable metric and data definition standard adopted across 6 client portfolios, contributing to governance and responsible AI documentation and explaining technical trade-offs to non-technical stakeholders.

Founder, Sharpmind, online tutoring agency — Jul 2026 - Present

- Built and deployed an LLM agent handling enquiries, scheduling and admin, automating roughly 70% of a real operational workload, then cut manual intervention by a further 60% through controlled A/B testing across 100+ live interactions.
- Teach maths, sciences and computer science to 25+ students in small groups, which is daily practice at coaching non-specialists through technical material.

Sales Manager, Flatpay, B2B fintech — Dec 2025 - Feb 2026

- Sold payment technology directly to business owners, reading an unfamiliar operation quickly and explaining a technical product and its trade-offs in plain commercial language.

Data Science Intern, Xceedure Business Solutions — Mar 2025 - Aug 2025

- Built machine learning models in Python with feature engineering and cross-validation using scikit-learn, XGBoost and LightGBM, and cleaned 30,000+ records of messy operational data causing 15-20% reporting inaccuracy.

PROJECTS

Agentic triage system with an assurance and evaluation layer

Stack: Python, LangGraph, FastAPI, E2B sandbox, Docker, OpenTelemetry, LangSmith

- Reached 88% severity classification accuracy across 25 labelled cases with a supervisor and specialist agent architecture that reproduces reported issues in a sandbox before classifying them.
- Built the reliability layer that makes it safe to run unattended: confidence-gated human escalation, schema validation, a dead letter queue, distributed tracing and an automated regression evaluation harness.

Retrieval and question answering pipeline, Natural Questions

Stack: Python, MiniLM dense retrieval, TF-IDF, reciprocal rank fusion, cross-encoder reranking

- Reached 93.80% Recall@1 across 871 queries with hybrid retrieval and cross-encoder reranking, against 54.76% for a TF-IDF baseline, and reported the failed query-expansion arm rather than dropping it.

Risk-aware allocation with deep reinforcement learning

Stack: Python, PyTorch, Gymnasium, PPO, MLflow, Docker, CI

- Took research code to deployable quality with zero constraint violations across 10,000 simulated decisions, sub-20ms p95 inference latency, automated leakage tests, MLflow experiment tracking, Dockerised inference and CI.

EDUCATION

BSc Mathematics, 2:1, Queen Mary University of London — Sept 2022 - May 2026

Relevant modules: Machine Learning; Natural Language Processing; Linear Algebra; Probability; Statistical Modelling.

TECHNICAL SKILLS

LLM Platforms and Generative AI: Anthropic Claude, OpenAI, retrieval-augmented generation, embeddings and vector retrieval, hybrid retrieval and reranking, structured outputs, tool calling, prompt and context engineering, fine-tuning (LoRA, RLHF)

Agent Design and Orchestration: LangGraph, LangChain, LangSmith, supervisor and specialist architectures, workflow automation, schema validation, confidence gating, evaluation harnesses, drift monitoring, OpenTelemetry

Cloud, Deployment and CI/CD: AWS, Azure, Databricks, Docker, Git, GitHub Actions CI/CD, MLflow, FastAPI, REST APIs, automated testing, Cursor and GitHub Copilot

Programming and ML: Python, SQL, C++, TypeScript, PyTorch, scikit-learn, XGBoost, LightGBM, reinforcement learning (PPO), feature engineering, cross-validation, model evaluation (precision, recall, F1, AUC)